

GIANT: Globally Improved Approximate Newton Method for Distributed Optimization

Shusen Wang

UC Berkeley

Joint work with **Fred Roosta, Peng Xu, and Michael Mahoney**


Background & Motivation

Optimization


- We consider the *empirical risk minimization* problem:

$$\min_{\mathbf{w} \in \mathbb{R}^d} \quad \left\{ f(\mathbf{w}) \triangleq \frac{1}{n} \sum_{j=1}^n l(\mathbf{w}; \mathbf{x}_j, y_j) + r(\mathbf{w}) \right\}$$

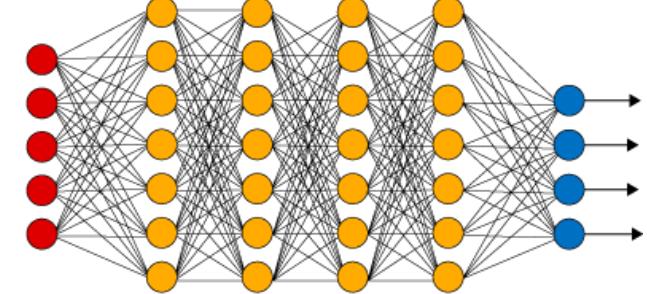
- Examples:



Linear Regression



Linear Classification



Neural Networks

Optimization

- How to solve the optimization problem $\min_w f(\mathbf{w})$?
 1. Write some code / find a package.




Optimization

- How to solve the optimization problem $\min_w f(\mathbf{w})$?
 1. Write some code / find a package.
 2. Load data to memory.



Disk



Random-access memory

Optimization

- How to solve the optimization problem $\min_w f(\mathbf{w})$?
 1. Write some code / find a package.
 2. Load data to memory.
 3. Run the code.



Optimization

- How to solve the optimization problem $\min_w f(\mathbf{w})$?
 1. Write some code / find a package.
 2. Load data to memory.
 3. Run the code.
- What if the data do not fit in memory?

Optimization

- How to solve the optimization problem $\min_w f(\mathbf{w})$?
 1. Write some code / find a package.
 2. Load data to memory.
 3. Run the code.
- What if the data do not fit in memory?
- What if the computation is too expensive for a single machine?

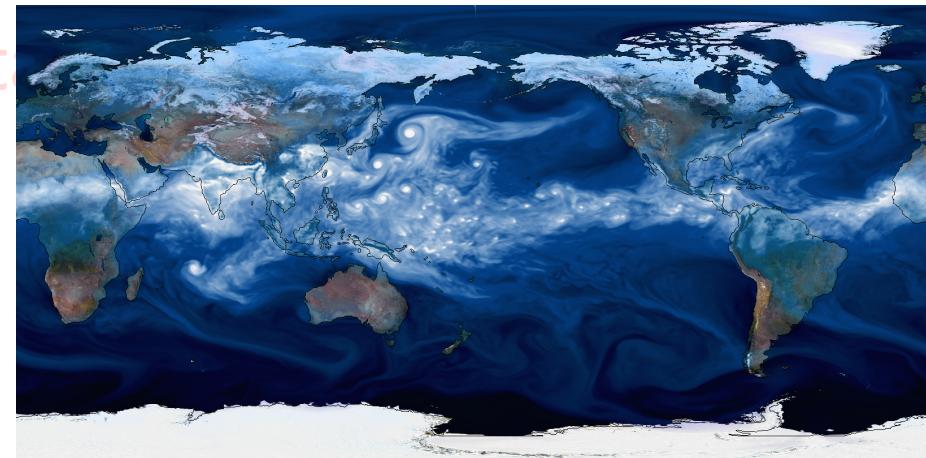
Optimization




3. Run the code.



- What if the data do not fit in memory?



Distributed Optimization



Computer clusters


Distributed Optimization

AMAZON EC2




Supercomputer

Distributed Optimization



Distributed Optimization



- $(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_n, y_n)$ are split among m machines.

Distributed Optimization

Ideally,

- $\frac{1}{m}$ of the data fit in the memory of one machine;
- each machine does $\frac{1}{m}$ of the computation $\rightarrow m$ x Speedup .

Distributed Optimization

Ideally,

- $\frac{1}{m}$ of the data fit in the memory of one machine;
- each machine does $\frac{1}{m}$ of the computation $\rightarrow \cancel{m \times \text{Speedup}}$.

Do not overlook the communication!

Distributed Optimization: Example

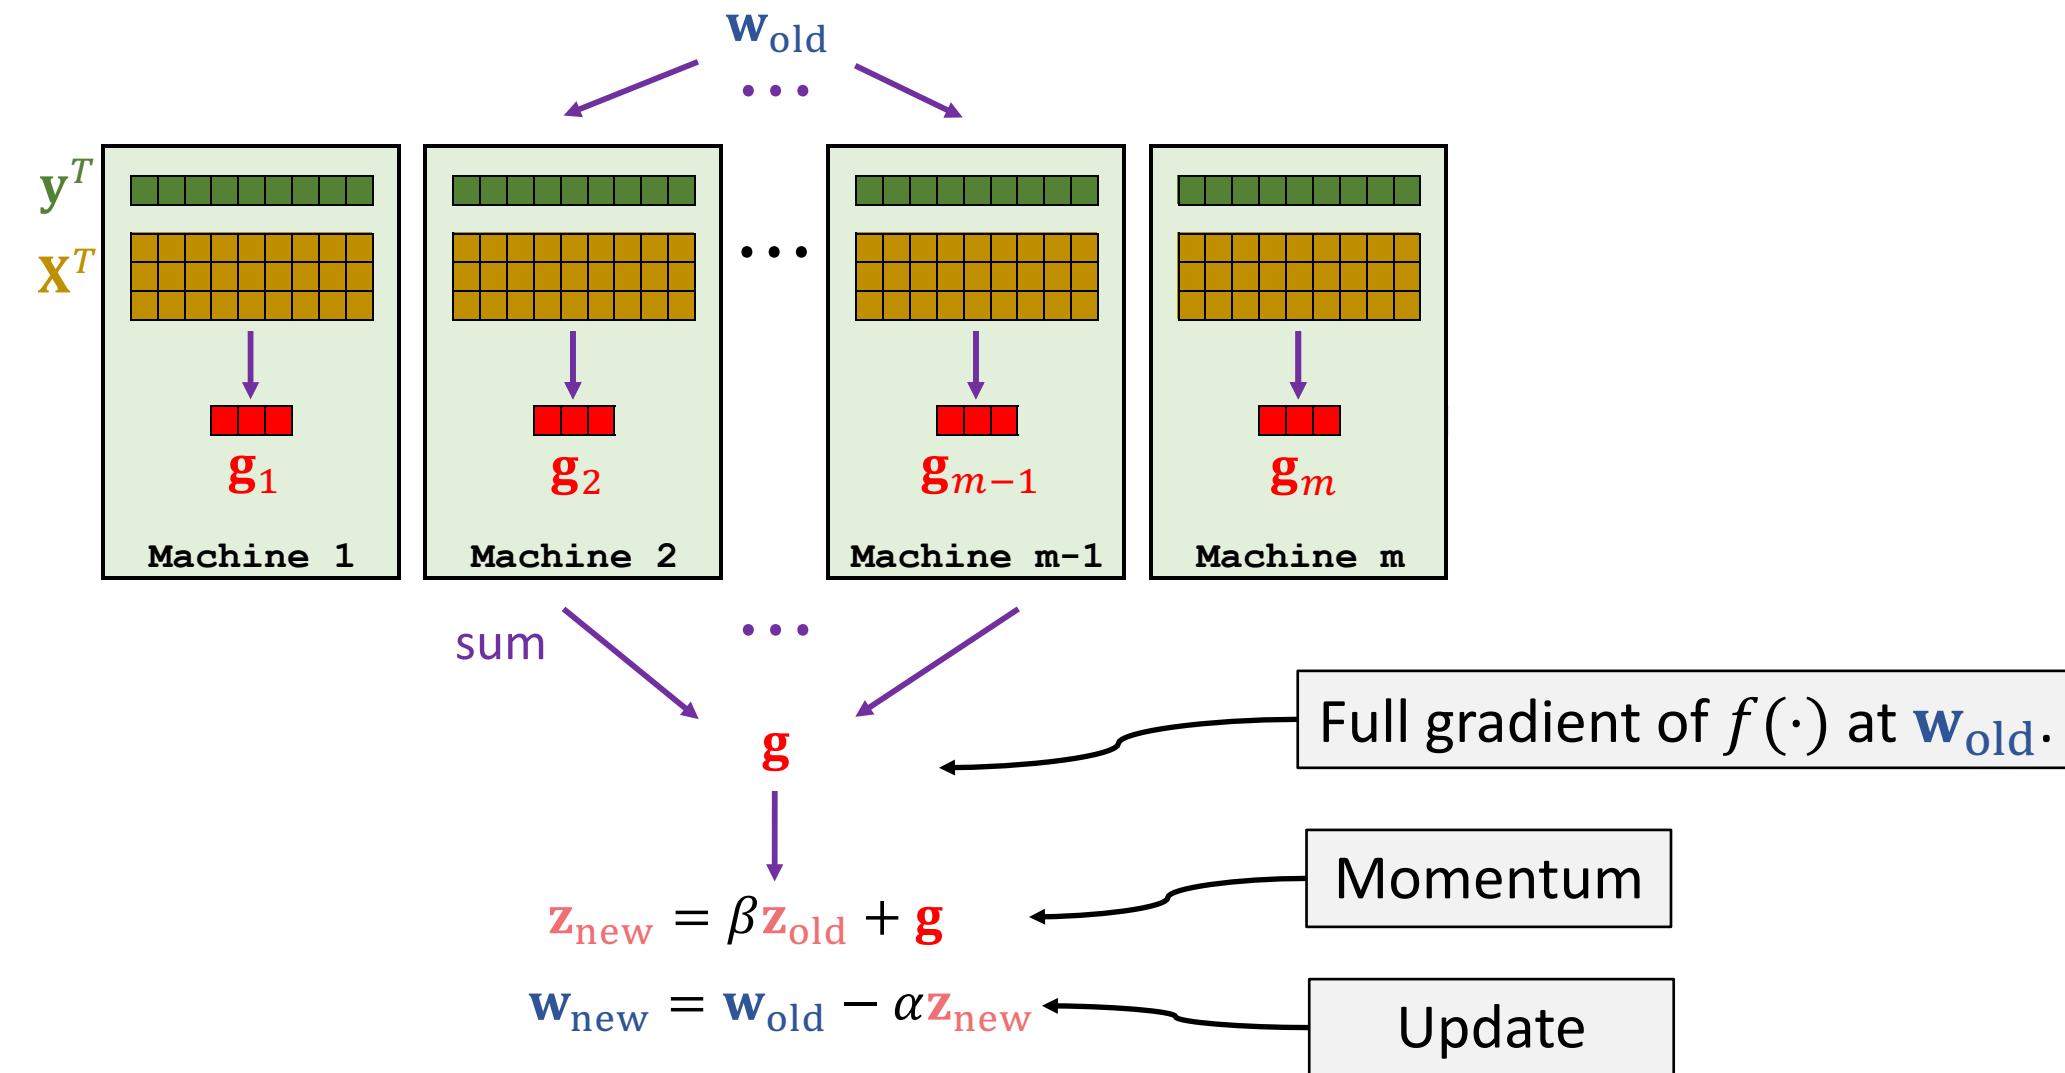
Solve the problem:

$$\min_{\mathbf{w} \in \mathbb{R}^d} \quad \left\{ f(\mathbf{w}) \triangleq \frac{1}{n} \sum_{j=1}^n l(\mathbf{w}; \mathbf{x}_j, y_j) + r(\mathbf{w}) \right\}$$

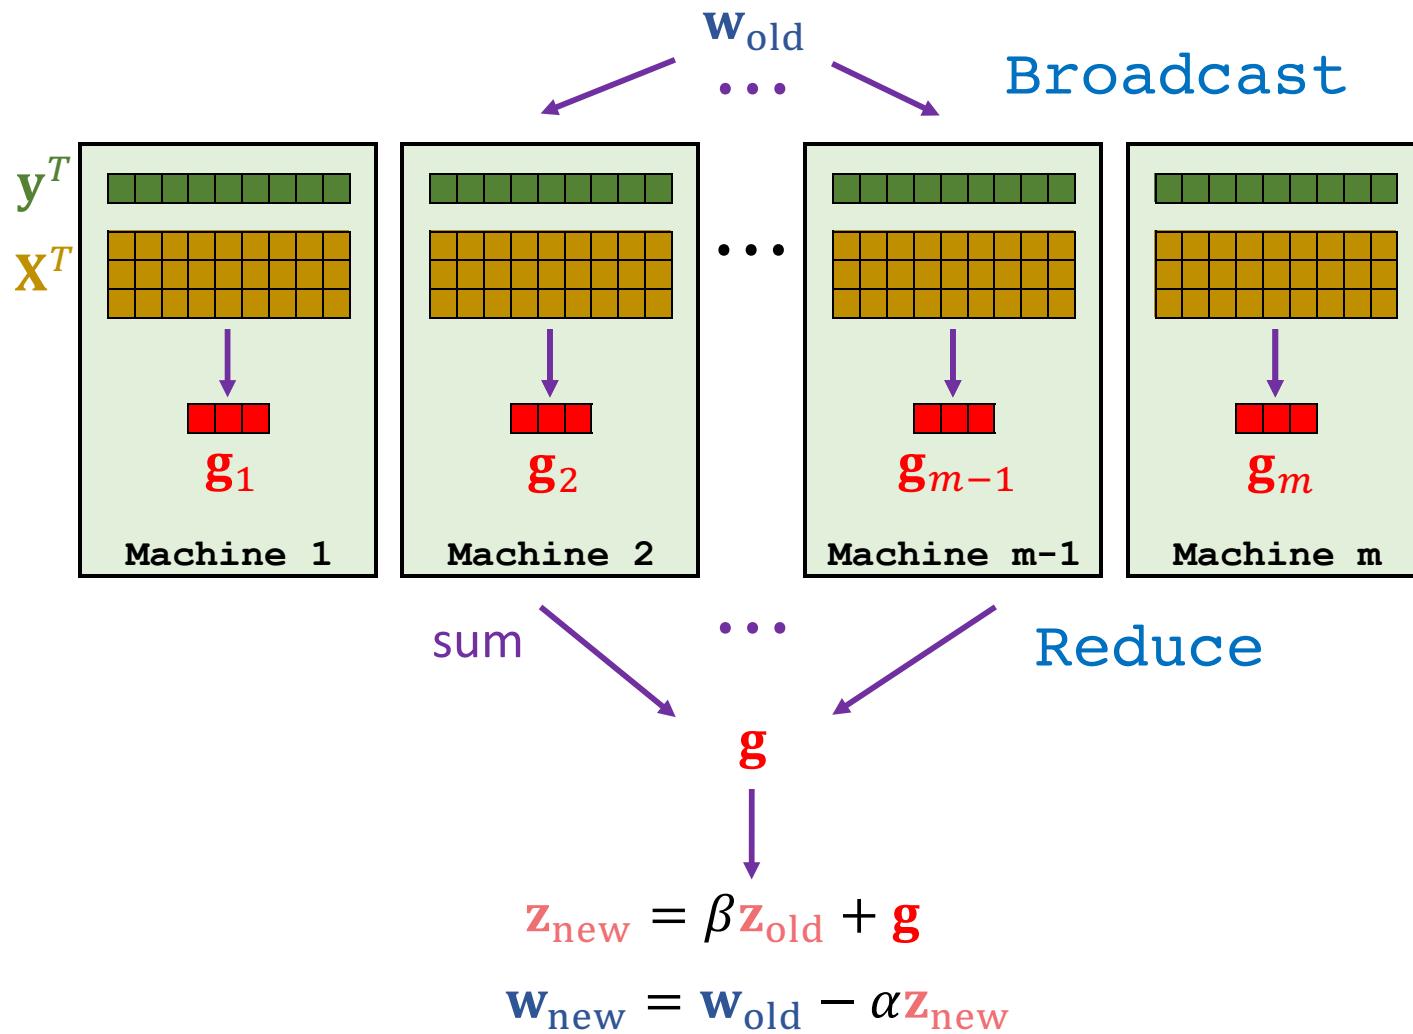
Accelerated Gradient Descent (AGD) repeats:

1. Compute gradient: $\mathbf{g} = \nabla f(\mathbf{w}_{\text{old}});$
2. Update momentum: $\mathbf{z}_{\text{new}} = \beta \mathbf{z}_{\text{old}} + \mathbf{g}, \quad 0 \leq \beta < 1;$
3. Update model: $\mathbf{w}_{\text{new}} = \mathbf{w}_{\text{old}} - \alpha \mathbf{z}_{\text{new}}.$

Warm-up: Distributed AGD

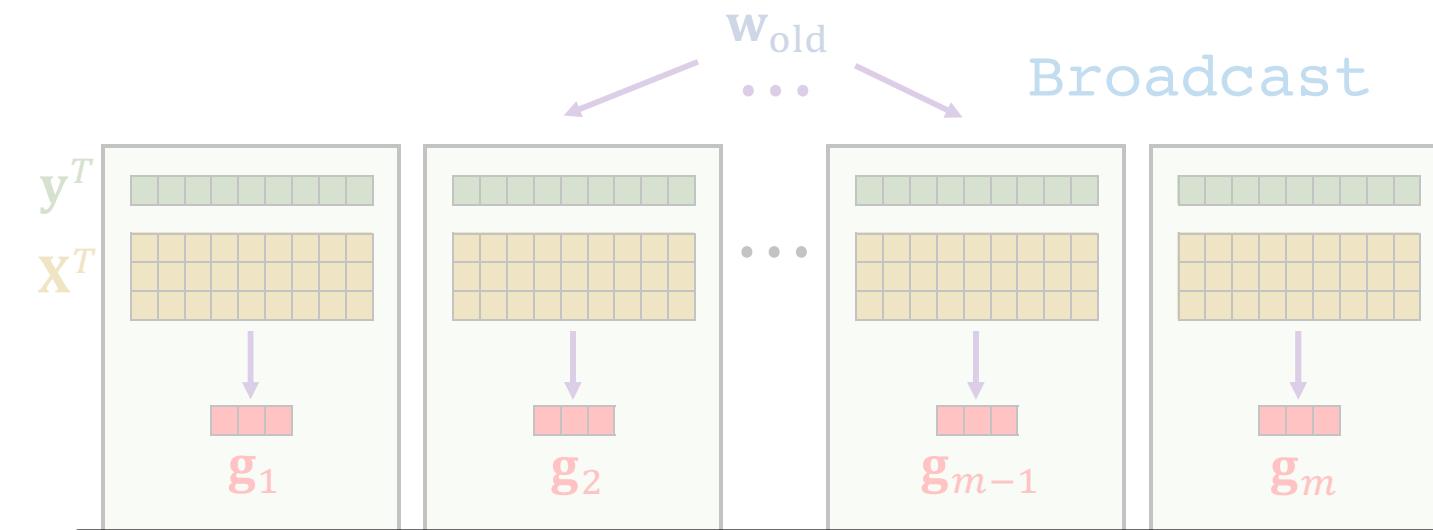


Warm-up: Distributed AGD



- Time complexity:
 $O\left(\frac{nd}{m}\right)$ FLOPs per iteration.
- One **Broadcast** and one **Reduce** per iteration.
- Lots of iterations to converge → lots of communications.

Warm-up: Distributed AGD



- Time complexity: $O\left(\frac{nd}{m}\right)$ FLOPs per iteration.
- One Broadcast and one Reduce

Cost = Computation + Communication

$$g \\ \downarrow \\ z_{\text{new}} = \beta z_{\text{old}} + g$$

$$w_{\text{new}} = w_{\text{old}} - \alpha z_{\text{new}}$$


- Lots of iterations to converge → lots of communications.

AGD for ℓ_2 -Regularized Logistic Regression

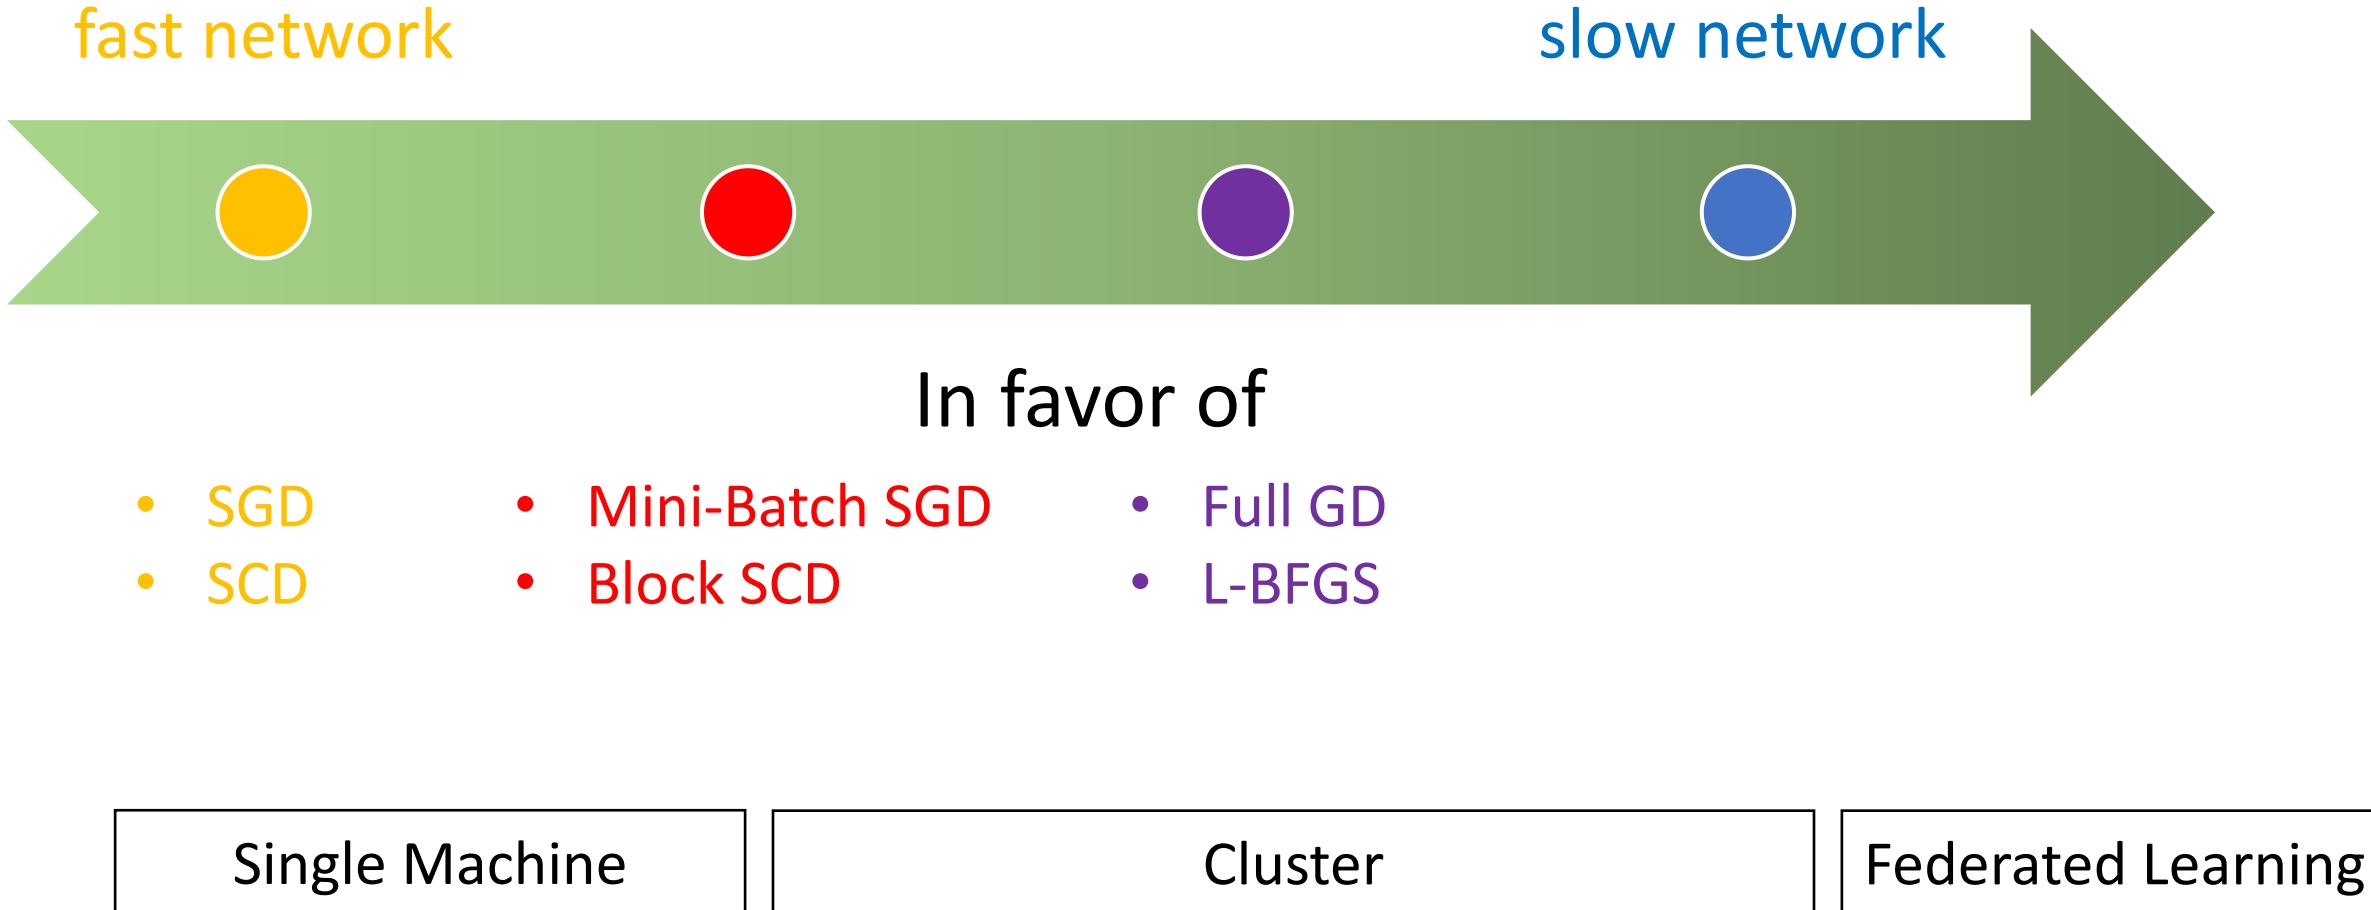
$$f(\mathbf{w}_t) - f(\mathbf{w}^*)$$

iteration

optimal
solution

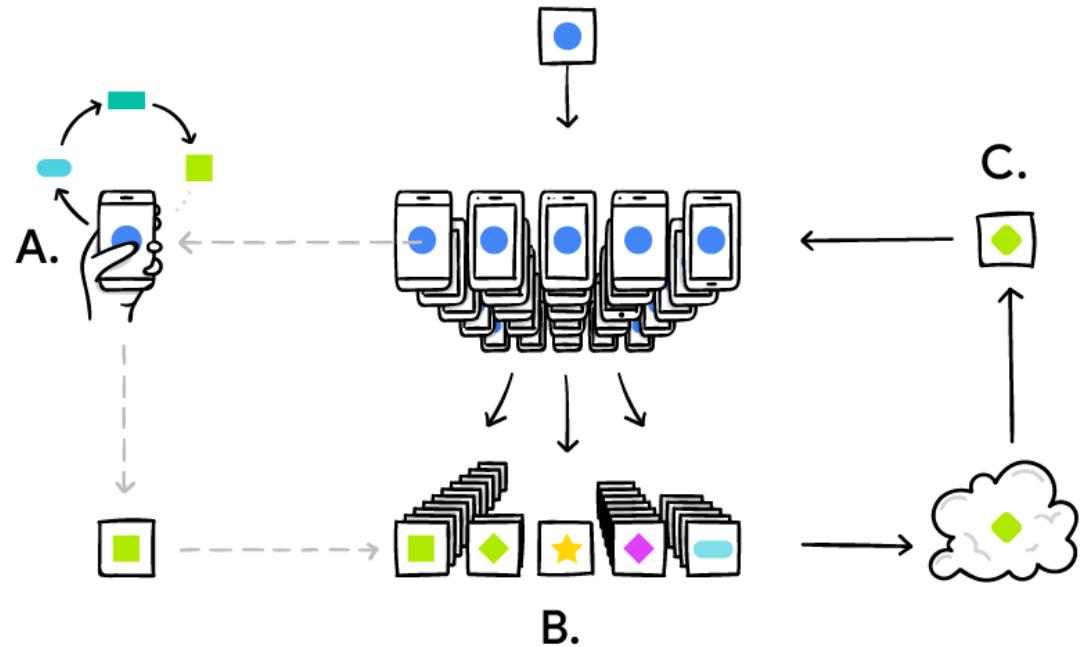


FLOPs versus Communications



FLOPs versus Communications

FLOPs versus Communications




The image is From Google Research Blog

Single Machine

Cluster

Federated Learning

FLOPs versus Communications



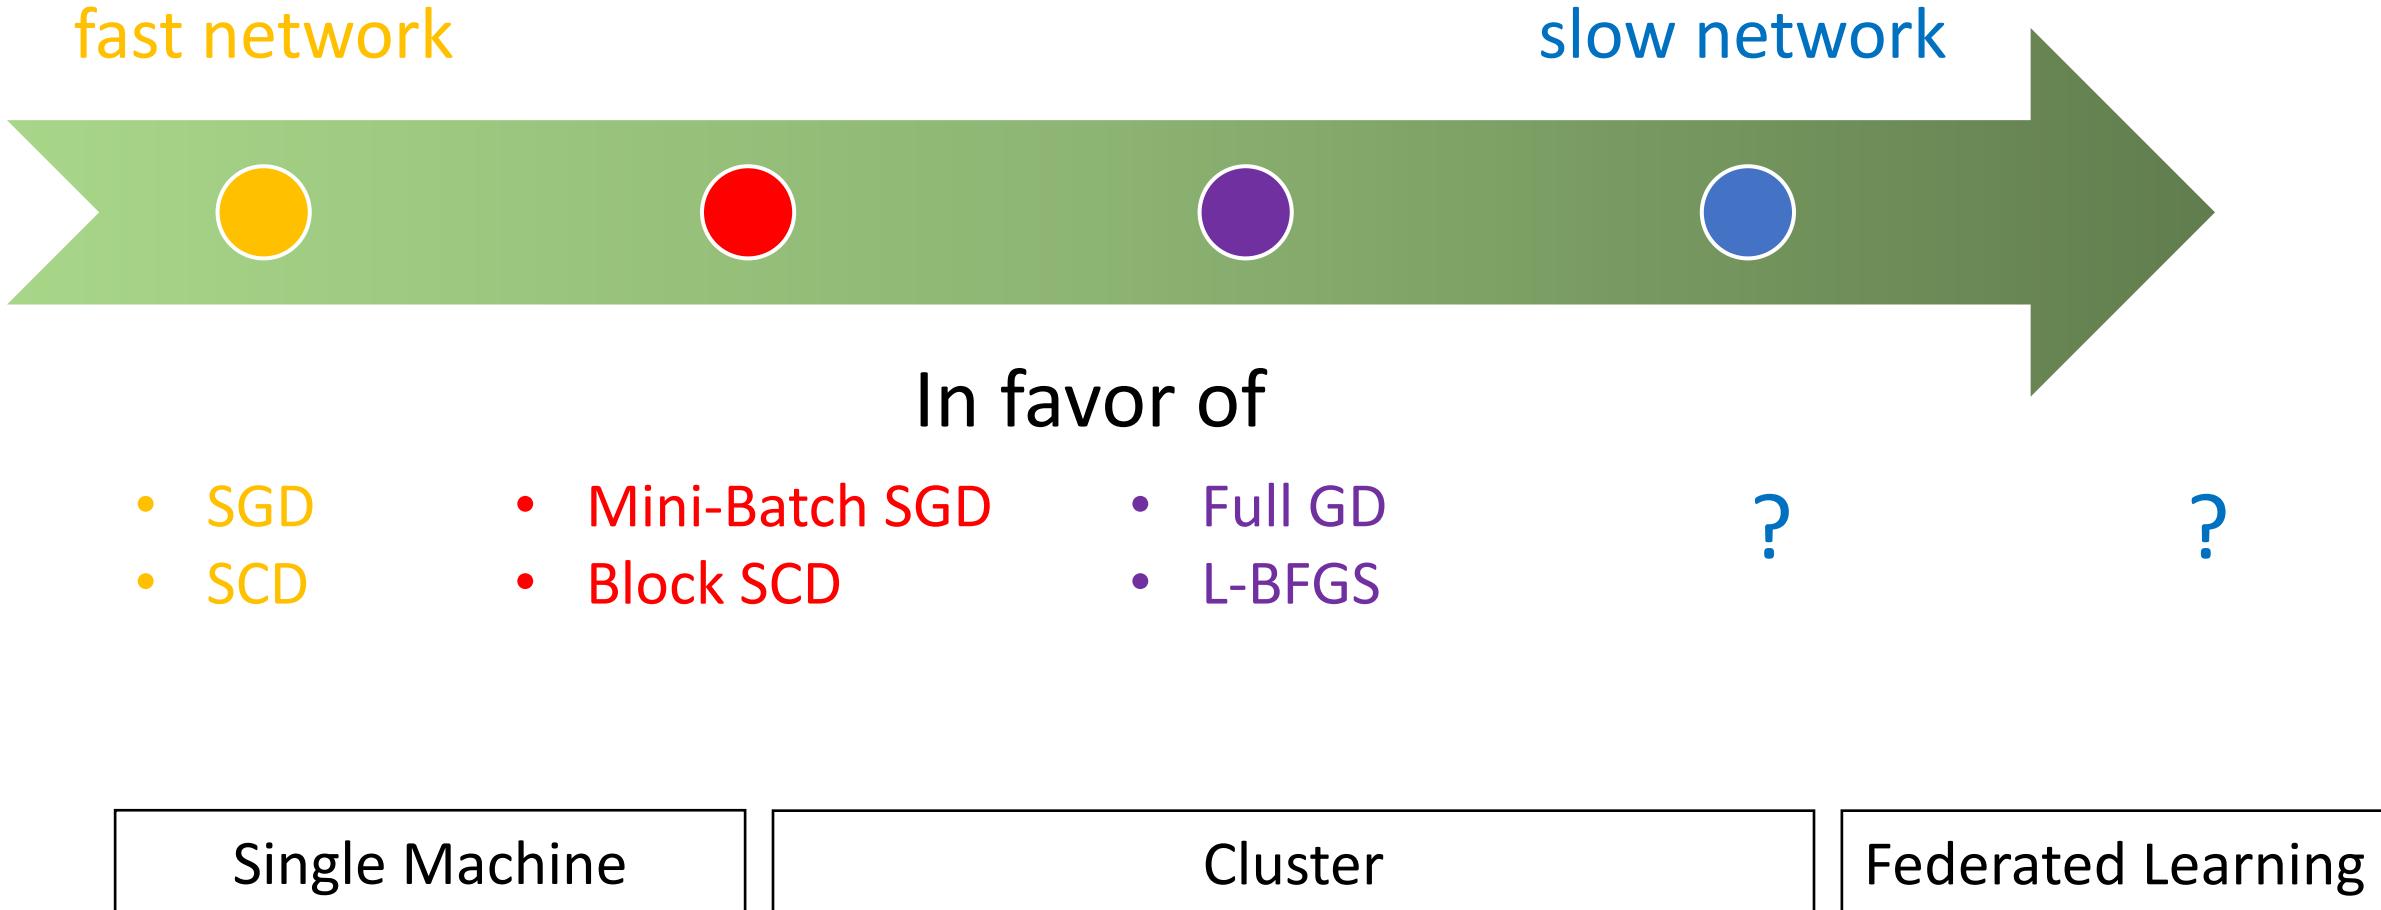
The image is From Google Research Blog

Single Machine

Cluster

Federated Learning

FLOPs versus Communications



Distributed Optimization

Summary

1. For big-data problems, distributed optimization is very useful.
2. If the network is slow, then communication is the bottleneck.
 - Recall: $\text{Cost} \approx \text{Computation} + \text{Communication}$

Communication-Efficient Optimization

Motivation

Basic ideas:

1. Let worker machines do lots of local computations.
2. Communicate as few as possible.

Prior Work

Existing communication-efficient methods:

- CoCoA
- DANE
- AIDE
-
-
-

They make assumptions, e.g.,

- objective function is **strongly convex** and **Lipschitz smooth**

Reference:

1. Smith, Forte, Ma, Takac, Jordan, & Jaggi. [CoCoA: A General Framework for Communication-Efficient Distributed Optimization](#).
2. Shamir, Srebro, & Zhang. [Communication Efficient Distributed Optimization using an Approximate Newton-type Method](#). In *ICML*, 2014.
3. Reddi, Konečný, Richtárik, Póczós, & Smola. [AIDE: Fast and Communication Efficient Distributed Optimization](#).

⋮

Prior Work

Existing communication-efficient methods:

- CoCoA
- DANE
- AIDE
-
-
-

Recall **Accelerated Gradient Descent (AGD)**

- $O\left(\sqrt{\kappa} \log \frac{1}{\epsilon}\right)$ iterations
- 2 communications per iteration
- $O\left(\frac{nd}{m}\right)$ FLOPs per iterations

Baseline!

Prior Work

Existing communication-efficient methods:

- CoCoA
- DANE
- AIDE
-
-
-

Recall **Accelerated Gradient Descent (AGD)**

- $O\left(\sqrt{\kappa} \log \frac{1}{\epsilon}\right)$ iterations
- 2 communications per iteration
- $O\left(\frac{nd}{m}\right)$ FLOPs per iterations

Baseline!

Do their convergence bounds beat **AGD**?

Prior Work

Existing communication-efficient methods:

- CoCoA
- DANE
- AIDE
-
-
-

Recall **Accelerated Gradient Descent (AGD)**

- $O\left(\sqrt{\kappa} \log \frac{1}{\epsilon}\right)$ iterations
- 2 communications per iteration
- $O\left(\frac{nd}{m}\right)$ FLOPs per iterations

Baseline!

Do their convergence bounds beat **AGD**?

- In terms of **communication**, NO!
- In terms of **computation**, NO!

Prior Work

Existing communication-efficient methods:

- CoCoA
- DANE
- AIDE

If the objective function is quadratic, then DANE = GIANT!

GIANT: Overview

Globally Improved Approximate Newton (GIANT)

- GIANT is a distributed 2nd-order method.
- Each iteration has 4 rounds of communications.
 - Broadcast or Reduce of one vector.
- Much faster convergence than AGD in terms of communication.
 - Assume the objective function is strongly convex and Lipschitz smooth.



Globally Improved Approximate Newton (**GIANT**)


- Assume the objective function is **strongly convex** and **Lipschitz smooth**.

Globally Improved Approximate Newton (GIANT)

- Assume the objective function is **strongly convex** and **Lipschitz smooth**.
- **Examples**


Linear regression

$$f(\mathbf{w}) = \frac{1}{2n} \sum_{j=1}^n (\mathbf{w}^T \mathbf{x}_j - y_j)^2 + \gamma \|\mathbf{w}\|_2^2$$



Logistic regression


$$f(\mathbf{w}) = \frac{1}{2n} \sum_{j=1}^n \log(1 + e^{-y_j \mathbf{w}^T \mathbf{x}_j}) + \gamma \|\mathbf{w}\|_2^2$$




Globally Improved Approximate Newton (GIANT)

- Assume the objective function is **strongly convex** and **Lipschitz smooth**.
- **Counter-examples**

LASSO



$$f(\mathbf{w}) = \frac{1}{2n} \sum_{j=1}^n (\mathbf{w}^T \mathbf{x}_j - y_j)^2 + \gamma \|\mathbf{w}\|_1$$


A graph showing the LASSO loss function. The x-axis is labeled x and ranges from -2 to 2. The y-axis is labeled $|x|$ and ranges from 0.5 to 2.0. The function is plotted as a red V-shape with a sharp corner at the origin (0,0).




Globally Improved Approximate Newton (GIANT)

- Assume the objective function is **strongly convex** and **Lipschitz smooth**.
- **Counter-examples**



Globally Improved Approximate Newton (GIANT)

- Assume the objective function is **strongly convex** and **Lipschitz smooth**.
- **Counter-examples**



Globally Improved Approximate Newton (**GIANT**)

- Assume the objective function is **strongly convex** and **Lipschitz smooth**.

GIANT: Algorithm Description

Warm-up: Newton-CG

- Repeat until convergence
 1. Compute gradient \mathbf{g} and Hessian \mathbf{H} ;
 2. Solve $\mathbf{H}\mathbf{p} = \mathbf{g}$ by running tens/hundreds of CG steps;
 3. Update $\mathbf{w} \leftarrow \mathbf{w} - \alpha\mathbf{p}$ (find α by line search).

GIANT: Algorithm Derivation

Recall: Newton's direction is $\mathbf{p} = \mathbf{H}^{-1}\mathbf{g}$.

In parallel, form the approximations:

$$\tilde{\mathbf{H}}_1 \approx \mathbf{H}$$

$$\tilde{\mathbf{H}}_2 \approx \mathbf{H}$$

...

$$\tilde{\mathbf{H}}_{m-1} \approx \mathbf{H}$$

$$\tilde{\mathbf{H}}_m \approx \mathbf{H}$$

In parallel, compute

$$\tilde{\mathbf{p}}_1 = \tilde{\mathbf{H}}_1^{-1}\mathbf{g}$$

$$\tilde{\mathbf{p}}_2 = \tilde{\mathbf{H}}_2^{-1}\mathbf{g}$$

...

$$\tilde{\mathbf{p}}_{m-1} = \tilde{\mathbf{H}}_{m-1}^{-1}\mathbf{g}$$

$$\tilde{\mathbf{p}}_m = \tilde{\mathbf{H}}_m^{-1}\mathbf{g}$$

$$\tilde{\mathbf{p}} = \frac{1}{m} \sum_i \tilde{\mathbf{p}}_i = \left(\frac{1}{m} \sum_i \tilde{\mathbf{H}}_i^{-1} \right) \mathbf{g}$$

approximates $\mathbf{p} = \mathbf{H}^{-1}\mathbf{g}$

GIANT: Algorithm Derivation

Recall: Newton's direction is $\mathbf{p} = \mathbf{H}^{-1}\mathbf{g}$.


$$\tilde{\mathbf{p}} = \frac{1}{m} \sum_i \tilde{\mathbf{p}}_i = \left(\frac{1}{m} \sum_i \tilde{\mathbf{H}}_i^{-1} \right) \mathbf{g} \quad \text{approximates } \mathbf{p} = \mathbf{H}^{-1}\mathbf{g}$$


GIANT: Algorithm Derivation


Recall: Newton's direction is $\mathbf{p} = \mathbf{H}^{-1}\mathbf{g}$.


$$\tilde{\mathbf{p}} = \frac{1}{m} \sum_i \tilde{\mathbf{p}}_i = \left(\frac{1}{m} \sum_i \tilde{\mathbf{H}}_i^{-1} \right) \mathbf{g} \quad \text{approximates } \mathbf{p} = \mathbf{H}^{-1}\mathbf{g}$$

- GIANT uses the exact gradient \mathbf{g} .
- GIANT approximates the Hessian matrix \mathbf{H} by $\left(\frac{1}{m} \sum_i \tilde{\mathbf{H}}_i^{-1} \right)^{-1}$.












most computations
are done here
(in parallel)




most computations
are done here
(in parallel)



most computations
are done here
(in parallel)



most computations
are done here
(in parallel)




Naïve approach:

1. Form local Hessian $\tilde{\mathbf{H}}_i \in \mathbb{R}^{d \times d}$
2. Invert $\tilde{\mathbf{H}}_i$
3. The ANT direction $\tilde{\mathbf{p}}_{t,i} = \tilde{\mathbf{H}}_i^{-1} \mathbf{g}_t$

It is inefficient!

1. Multiply two matrices to form $\tilde{\mathbf{H}}_i$
2. Invert the dense matrix $\tilde{\mathbf{H}}_i$


most computations
are done here
(in parallel)




Fact: For the problem

$$\min_{\mathbf{w} \in \mathbb{R}^d} \left\{ f(\mathbf{w}) \triangleq \frac{1}{n} \sum_{j=1}^n l(\mathbf{w}; \mathbf{x}_j, y_j) + \gamma \|\mathbf{w}\|_2^2 \right\},$$

the local Hessian can be written as $\tilde{\mathbf{H}}_i = \mathbf{A}_i^T \mathbf{A}_i + \gamma \mathbf{I}_d$.



most computations
are done here
(in parallel)



Fact: For the problem


$$\min_{\mathbf{w} \in \mathbb{R}^d} \left\{ f(\mathbf{w}) \triangleq \frac{1}{n} \sum_{j=1}^n l(\mathbf{w}; \mathbf{x}_j, y_j) + \gamma \|\mathbf{w}\|_2^2 \right\},$$


the local Hessian can be written as $\tilde{\mathbf{H}}_i = \mathbf{A}_i^T \mathbf{A}_i + \gamma \mathbf{I}_d$.

Local solver:

- Inexactly solve $(\mathbf{A}_i^T \mathbf{A}_i + \gamma \mathbf{I}_d) \mathbf{p} = \mathbf{g}_t$ by taking q CG steps.
- Cost: $2q$ matrix-vector products.

most computations
are done here
(in parallel)





GIANT: Experiments

Settings

- Solve the ℓ_2 -regularized logistic regression:


$$\min_{\mathbf{w} \in \mathbb{R}^d} \quad \left\{ f(\mathbf{w}) \triangleq \frac{1}{n} \sum_{j=1}^n \log \left(1 + e^{-y_j \mathbf{x}_j^T \mathbf{w}} \right) + \frac{\gamma}{2} \|\mathbf{w}\|_2^2 \right\}$$

Datasets

- Covtype: $n = 581K, d = 54$.
- Epsilon: $n = 500K, d = 2K$.
- 80% for training, 20% for test.

Datasets

- Covtype: $n = 581K, d = 54$.
- Epsilon: $n = 500K, d = 2K$.
- 80% for training, 20% for test.



Compared Methods

- Accelerated gradient descent (AGD)
 - choose *step size* from {0.1, 1, 10, 100}
 - choose *momentum* from {0.5, 0.9, 0.95, 0.99, 0.999}

Compared Methods

- Accelerated gradient descent (AGD)
- Limited memory BFGS (a quasi-Newton method)
 - choose *number of history* from {30, 100, 300}
 - line search is used

Compared Methods

- Accelerated gradient descent (AGD)
- Limited memory BFGS
- DANE (another Newton-type method) [Shamir et al. 2014]
 - local solver: **SVRG** (a stochastic optimization method)
 - choose *step size of SVRG* from {0.1, 1, 10, 100}
 - choose *max. iteration of SVRG* from {30, 100, 300}

Reference:

Shamir, Srebro, & Zhang. Communication Efficient Distributed Optimization using an Approximate Newton-type Method. In *ICML*, 2014.

Compared Methods

- Accelerated gradient descent (AGD)
- Limited memory BFGS
- DANE (another Newton-type method)
- GIANT
 - local solver: conjugate gradient (CG)
 - choose *max iteration of CG* from {30, 100, 300}

Compared Methods

- Accelerated gradient descent (AGD)
- Limited memory BFGS
- DANE (another Newton-type method)
- GIANT

2 Tuning Parameters

1 Tuning Parameter

2 Tuning Parameters

1 Tuning Parameter

Experiment Environment

- Spark 2.1.1 + Scala 2.11.8



Experiment Environment

- Spark 2.1.1 + Scala 2.11.8
- Cori Supercomputer (Cray XC40)




National Energy Research
Scientific Computing Center




Experiment Environment

- Spark 2.1.1 + Scala 2.11.8
- Cori Supercomputer (Cray XC40)
 - 128 GB Memory / node
 - 32 Cores / node
- Use 15 nodes (480 CPU cores)


Covtype (n=581K, $\bar{d}=10K$), Training




Epsilon ($n=500K$, $\bar{d}=10K$), Training



Covtype (n=581K, $\bar{d}=10K$), Test




Epsilon ($n=500K$, $\bar{d}=10K$), Test




Scaling Experiments

- Make the Covtype data k times larger.
 1. Get k replicates of \mathbf{X} and \mathbf{y} ;
 2. Inject i.i.d. Gaussian noises to the $kn \times d$ feature matrix;
 3. Do random feature mapping to get 10K features.
- Use k times more nodes.
- Set $k = 5$ and $k = 25$.


Original Data, 15 Nodes (480 Cores)



5x Larger Data, 75 Nodes (2.4K Cores)



25x Larger Data, 375 Nodes (12K Cores)





Why is GIANT More Scalable?

- As **#Samples** and **#Nodes** both increases by k times,
 - the **computational** costs remain **the same**;
 - the **communication** costs **increase**.

Why is GIANT More Scalable?


- As #Samples and #Nodes both increases by k times,
 - the computational costs remain the same;
 - the communication costs increase.
- Per-iteration time of AGD and L-BFGS increases.



Why is GIANT More Scalable?

- As #Samples and #Nodes both increases by k times,
 - the computational costs remain the same;
 - the communication costs increase.
- Per-iteration time of AGD and L-BFGS increases.
- Per-iteration time of GIANT marginally increases.
 - Because GIANT is computation-intensive.

FLOPs versus Communication




GIANT: Convergence Analysis

Quadratic Loss: Global Convergence

- Objective function: $f(\mathbf{w}) = \frac{1}{2n} \|\mathbf{X}\mathbf{w} - \mathbf{y}\|_2^2 + \frac{\gamma}{2} \|\mathbf{w}\|_2^2$

Quadratic Loss: Global Convergence


- Objective function: $f(\mathbf{w}) = \frac{1}{2n} \|\mathbf{X}\mathbf{w} - \mathbf{y}\|_2^2 + \frac{\gamma}{2} \|\mathbf{w}\|_2^2$
 - Assume \mathbf{X} is “incoherent” (information uniformly spread)
 - Assume local sample size is $s = \Theta\left(\frac{d}{\epsilon^2} \log \frac{md}{\delta}\right)$ for any $\epsilon, \delta \in (0, 1)$



Quadratic Loss: Global Convergence

- Objective function: $f(\mathbf{w}) = \frac{1}{2n} \|\mathbf{X}\mathbf{w} - \mathbf{y}\|_2^2 + \frac{\gamma}{2} \|\mathbf{w}\|_2^2$
- Assume \mathbf{X} is “incoherent” (information uniformly spread)
- Assume local sample size is $s = \Theta\left(\frac{d}{\epsilon^2} \log \frac{md}{\delta}\right)$ for any $\epsilon, \delta \in (0, 1)$
- With probability $1 - \delta$ (assume random partition of the data),


$$\frac{\|\Delta_t\|_2}{\|\Delta_0\|_2} \leq \left(\frac{\epsilon}{\sqrt{m}} + \epsilon^2\right)^t \sqrt{\kappa}, \quad \text{where } \Delta_t \triangleq \mathbf{w}_t - \mathbf{w}^*.$$



Quadratic Loss: Global Convergence


- Objective function: $f(\mathbf{w}) = \frac{1}{2n} \|\mathbf{X}\mathbf{w} - \mathbf{y}\|_2^2 + \frac{\gamma}{2} \|\mathbf{w}\|_2^2$
- Ass GIANT has $\log \kappa$ dependence.
- Assume local sample size is $s = \Theta\left(\frac{\alpha}{\epsilon^2} \log \frac{n\alpha}{\delta}\right)$ for any $\epsilon, \delta \in (0, 1)$
- With probability $1 - \delta$ (assume random partition of the data).

AGD has $\sqrt{\kappa}$ dependence.




General Smooth Loss: Local Convergence

- Denote $\mathbf{H}_t = \nabla^2 f(\mathbf{w}_t)$ and $\mathbf{H}^* = \nabla^2 f(\mathbf{w}^*)$
- Similar “incoherent” assumption (information uniformly spread)
- Assume local sample size is $s = \Theta\left(\frac{d}{\epsilon^2} \log \frac{md}{\delta}\right)$ for any $\epsilon, \delta \in (0, 1)$



General Smooth Loss: Local Convergence


- Denote $\mathbf{H}_t = \nabla^2 f(\mathbf{w}_t)$ and $\mathbf{H}^* = \nabla^2 f(\mathbf{w}^*)$
- Similar “incoherent” assumption (information uniformly spread)
- Assume local sample size is $s = \Theta\left(\frac{d}{\epsilon^2} \log \frac{md}{\delta}\right)$ for any $\epsilon, \delta \in (0, 1)$
- Assume the Hessian is L -Lipchitz: $\|\nabla^2 f(\mathbf{w}) - \nabla^2 f(\mathbf{w}')\|_2 \leq L\|\mathbf{w} - \mathbf{w}'\|_2$



General Smooth Loss: Local Convergence

- Denote $\mathbf{H}_t = \nabla^2 f(\mathbf{w}_t)$ and $\mathbf{H}^\star = \nabla^2 f(\mathbf{w}^\star)$
- Similar “incoherent” assumption (information uniformly spread)
- Assume local sample size is $s = \Theta\left(\frac{d}{\epsilon^2} \log \frac{md}{\delta}\right)$ for any $\epsilon, \delta \in (0, 1)$
- Assume the Hessian is L -Lipchitz: $\|\nabla^2 f(\mathbf{w}) - \nabla^2 f(\mathbf{w}')\|_2 \leq L\|\mathbf{w} - \mathbf{w}'\|_2$
- With probability $1 - \delta$ (assume random partition of the data),

$$\|\Delta_{t+1}\|_2 \leq \max \left\{ \left(\frac{\epsilon}{\sqrt{m}} + \epsilon^2 \right) \sqrt{\frac{\sigma_{\max}(\mathbf{H}_t)}{\sigma_{\min}(\mathbf{H}_t)}} \|\Delta_t\|_2, \frac{2L}{\sigma_{\min}(\mathbf{H}_t)} \|\Delta_t\|_2^2 \right\}$$



Inexactly Solving Local Linear System

- Exactly solving $\tilde{\mathbf{H}}_{t,i} \mathbf{p} = \mathbf{g}_t$ may not be easy.
- Solve $\tilde{\mathbf{H}}_{t,i} \mathbf{p} = \mathbf{g}_t$ by taking $q = \frac{\sqrt{\kappa}-1}{2} \log \frac{8}{\epsilon_0^2}$ CG steps.
- Recall the bounds of **exact** solver:

$$\text{Quadratic Loss: } \frac{\|\Delta_t\|_2}{\|\Delta_0\|_2} \leq \left(\frac{\epsilon}{\sqrt{m}} + \epsilon^2 \right)^t \sqrt{\kappa}, \quad \text{where } \Delta_t \triangleq \mathbf{w}_t - \mathbf{w}^*.$$

$$\text{General Loss: } \|\Delta_{t+1}\|_2 \leq \max \left\{ \left(\frac{\epsilon}{\sqrt{m}} + \epsilon^2 \right) \sqrt{\frac{\sigma_{\max}(\mathbf{H}_t)}{\sigma_{\min}(\mathbf{H}_t)}} \|\Delta_t\|_2, \frac{2L}{\sigma_{\min}(\mathbf{H}_t)} \|\Delta_t\|_2^2 \right\}$$

Inexactly Solving Local Linear System

- Exactly solving $\tilde{\mathbf{H}}_{t,i} \mathbf{p} = \mathbf{g}_t$ may not be easy.
- Solve $\tilde{\mathbf{H}}_{t,i} \mathbf{p} = \mathbf{g}_t$ by taking $q = \frac{\sqrt{\kappa}-1}{2} \log \frac{8}{\epsilon_0^2}$ CG steps.
- Recall the bounds of **exact** solver:

Quadratic Loss:
$$\frac{\|\Delta_t\|_2}{\|\Delta_0\|_2} \leq \left(\frac{\epsilon}{\sqrt{m}} + \epsilon^2 \right)^t \sqrt{\kappa}, \quad \text{where } \Delta_t \triangleq \mathbf{w}_t - \mathbf{w}^*.$$

General Loss:
$$\|\Delta_{t+1}\|_2 \leq \max \left\{ \left(\frac{\epsilon}{\sqrt{m}} + \epsilon^2 \right) \sqrt{\frac{\sigma_{\max}(\mathbf{H}_t)}{\sigma_{\min}(\mathbf{H}_t)}} \|\Delta_t\|_2, \frac{2L}{\sigma_{\min}(\mathbf{H}_t)} \|\Delta_t\|_2^2 \right\}$$

- Bounds of the inexact solver:

$$\frac{\epsilon}{\sqrt{m}} + \epsilon^2 \implies \frac{\epsilon}{\sqrt{m}} + \epsilon^2 + \epsilon_0$$

Inexactly Solving Local Linear System

- Exactly solving $\tilde{\mathbf{H}}_{t,i} \mathbf{p} = \mathbf{g}_t$ may not be easy.
- Solve $\tilde{\mathbf{H}}_{t,i} \mathbf{p} = \mathbf{g}_t$ by taking $q = \frac{\sqrt{\kappa}-1}{2} \log \frac{8}{\epsilon_0^2}$ CG steps.
- Recall the bounds of **exact** solver:

Quadratic Loss: $\frac{\|\Delta_t\|_2}{\|\Delta_0\|_2} \leq \left(\frac{\epsilon}{\sqrt{m}} + \epsilon^2 \right)^t \sqrt{\kappa}, \quad \text{where } \Delta_t \triangleq \mathbf{w}_t - \mathbf{w}^*.$

General Loss: $\|\Delta_{t+1}\|_2 \leq \max \left\{ \left(\frac{\epsilon}{\sqrt{m}} + \epsilon^2 \right) \sqrt{\frac{\sigma_{\max}(\mathbf{H}_t)}{\sigma_{\min}(\mathbf{H}_t)}} \|\Delta_t\|_2, \frac{2L}{\sigma_{\min}(\mathbf{H}_t)} \|\Delta_t\|_2^2 \right\}$

- Bounds of the inexact solver:

$$\frac{\epsilon}{\sqrt{m}} + \epsilon^2 \implies \frac{\epsilon}{\sqrt{m}} + \epsilon^2 + \boxed{\epsilon_0}$$

Outline of Proof


Proof Techniques

Claim 1: Local Hessian $\tilde{\mathbf{H}}_1, \dots, \tilde{\mathbf{H}}_m$ well approximate the true Hessian \mathbf{H} .

Proof Techniques

Claim 1: Local Hessian $\tilde{\mathbf{H}}_1, \dots, \tilde{\mathbf{H}}_m$ well approximate the true Hessian \mathbf{H} .


- \mathbf{H} can always be written as $\mathbf{H} = \mathbf{A}^T \mathbf{A} + \gamma \mathbf{I}_d$, for some $\mathbf{A} \in \mathbb{R}^{n \times d}$.



Proof Techniques

Claim 1: Local Hessian $\tilde{\mathbf{H}}_1, \dots, \tilde{\mathbf{H}}_m$ well approximate the true Hessian \mathbf{H} .

- \mathbf{H} can always be written as $\mathbf{H} = \mathbf{A}^T \mathbf{A} + \gamma \mathbf{I}_d$, for some $\mathbf{A} \in \mathbb{R}^{n \times d}$.
- Assume \mathbf{A} is “incoherent”.
- Uniformly partition \mathbf{A} to form $\mathbf{A}_i \in \mathbb{R}^{s \times d}$.



Proof Techniques

Claim 1: Local Hessian $\tilde{\mathbf{H}}_1, \dots, \tilde{\mathbf{H}}_m$ well approximate the true Hessian \mathbf{H} .

- \mathbf{H} can always be written as $\mathbf{H} = \mathbf{A}^T \mathbf{A} + \gamma \mathbf{I}_d$, for some $\mathbf{A} \in \mathbb{R}^{n \times d}$.
- Assume \mathbf{A} is “incoherent”.
- Uniformly partition \mathbf{A} to form $\mathbf{A}_i \in \mathbb{R}^{s \times d}$.

$$\begin{array}{c} \text{---} \\ \text{---} \\ \text{---} \\ \text{---} \\ \text{---} \end{array} \bullet \begin{array}{c} \text{---} \\ \text{---} \\ \text{---} \\ \text{---} \\ \text{---} \end{array} \approx \frac{n}{s} \bullet \begin{array}{c} \text{---} \\ \text{---} \\ \text{---} \\ \text{---} \\ \text{---} \end{array} \bullet \begin{array}{c} \text{---} \\ \text{---} \\ \text{---} \\ \text{---} \\ \text{---} \end{array}$$

$\mathbf{A}^T \qquad \mathbf{A} \qquad \qquad \qquad \mathbf{A}_i^T \qquad \mathbf{A}_i$

Proof Techniques

Claim 1: Local Hessian $\tilde{\mathbf{H}}_1, \dots, \tilde{\mathbf{H}}_m$ well approximate the true Hessian \mathbf{H} .

- \mathbf{H} can always be written as $\mathbf{H} = \mathbf{A}^T \mathbf{A} + \gamma \mathbf{I}_d$, for some $\mathbf{A} \in \mathbb{R}^{n \times d}$.
- Assume \mathbf{A} is “incoherent”.
- Uniformly partition \mathbf{A} to form $\mathbf{A}_i \in \mathbb{R}^{s \times d}$.
- Sufficiently large samples size $s = \Theta\left(\frac{d}{\epsilon^2} \log \frac{d}{\delta}\right)$
- By matrix Bernstein (concentration inequality), with probability $1 - \delta$,

$$(1 - \epsilon) \mathbf{A}^T \mathbf{A} \leq \frac{n}{s} \mathbf{A}_i^T \mathbf{A}_i \leq (1 + \epsilon) \mathbf{A}^T \mathbf{A}.$$

Proof Techniques

Claim 1: Local Hessian $\tilde{\mathbf{H}}_1, \dots, \tilde{\mathbf{H}}_m$ well approximate the true Hessian \mathbf{H} .

- \mathbf{H} can always be written as $\mathbf{H} = \mathbf{A}^T \mathbf{A} + \gamma \mathbf{I}_d$, for some $\mathbf{A} \in \mathbb{R}^{n \times d}$.
- Assume \mathbf{A} is “incoherent”.
- Uniformly partition \mathbf{A} to form $\mathbf{A}_i \in \mathbb{R}^{s \times d}$.
- Sufficiently large samples size $s = \Theta\left(\frac{d}{\epsilon^2} \log \frac{d}{\delta}\right)$
- By matrix Bernstein (concentration inequality), with probability $1 - \delta$,


$$(1 - \epsilon) \mathbf{A}^T \mathbf{A} \leq \frac{n}{s} \mathbf{A}_i^T \mathbf{A}_i \leq (1 + \epsilon) \mathbf{A}^T \mathbf{A}.$$

- Note that $\tilde{\mathbf{H}}_i = \frac{n}{s} \mathbf{A}_i^T \mathbf{A}_i + \gamma \mathbf{I}_d \longrightarrow \tilde{\mathbf{H}}_i$ well approximates \mathbf{H} .

Proof Techniques

Claim 2: The GIANT direction approximates $\mathbf{p}^* = \mathbf{H}^{-1}\mathbf{g}$.

- Define the quadratic function $\phi(\mathbf{p}) \triangleq \frac{1}{2}\mathbf{p}^T \mathbf{H} \mathbf{p} - \mathbf{p}^T \mathbf{g} \quad (\leq 0)$



Proof Techniques

Claim 2: The GIANT direction approximates $\mathbf{p}^* = \mathbf{H}^{-1}\mathbf{g}$.

- Define the quadratic function $\phi(\mathbf{p}) \triangleq \frac{1}{2}\mathbf{p}^T \mathbf{H} \mathbf{p} - \mathbf{p}^T \mathbf{g} \quad (\leq 0)$
- The exact Newton direction is $\mathbf{p}^* = \mathbf{H}^{-1}\mathbf{g} = \underset{\mathbf{p}}{\operatorname{argmin}} \phi(\mathbf{p})$

Proof Techniques

Claim 2: The GIANT direction approximates $\mathbf{p}^* = \mathbf{H}^{-1}\mathbf{g}$.

- Define the quadratic function $\phi(\mathbf{p}) \triangleq \frac{1}{2}\mathbf{p}^T \mathbf{H} \mathbf{p} - \mathbf{p}^T \mathbf{g} \quad (\leq 0)$
- The exact Newton direction is $\mathbf{p}^* = \mathbf{H}^{-1}\mathbf{g} = \underset{\mathbf{p}}{\operatorname{argmin}} \phi(\mathbf{p})$
- The GIANT directions is $\tilde{\mathbf{p}} \triangleq \frac{1}{m} \sum_{i=1}^m \tilde{\mathbf{p}}_i \triangleq \frac{1}{m} \sum_{i=1}^m \tilde{\mathbf{H}}_i^{-1}\mathbf{g}$
- Conditioning on **Claim 1** that $\tilde{\mathbf{H}}_i$ well approximates \mathbf{H} , we get

$$\phi(\mathbf{p}^*) \leq \phi(\tilde{\mathbf{p}}) \leq (1 - \alpha^2) \cdot \phi(\mathbf{p}^*), \quad \text{where } \alpha = \left(\frac{\epsilon}{\sqrt{m}} + \epsilon^2 \right)$$

Reference:

W, Gittens, & Mahoney: Sketched Ridge Regression: Optimization Perspective, Statistical Perspective, and Model Averaging. In *ICML* 2017.

Proof Techniques

1. Use **Claim 2** that $\phi(\mathbf{p}^*) \leq \phi(\tilde{\mathbf{p}}) \leq (1 - \alpha^2) \cdot \phi(\mathbf{p}^*)$, where $\alpha = (\frac{\epsilon}{\sqrt{m}} + \epsilon^2)$
2. Follow the standard convergence analysis of Newton's method.

→ Convergence of GIANT!

Conclusions & Future Work

Conclusions & Future Work

- GIANT's theory beats the existing works.
 - Assume the objective function is strongly convex and Lipschitz smooth.
- GIANT has good empirical performance on computer cluster.
 - Beats AGD, L-BFGS, and DANE.

Conclusions & Future Work

- GIANT's theory beats the existing works.
 - Assume the objective function is **strongly convex** and **Lipschitz smooth**.
- GIANT has good empirical performance on computer cluster.


Conclusions & Future Work

- GIANT's theory beats the existing works.
 - Assume the objective function is **strongly convex** and **Lipschitz smooth**.


Counter-examples

LASSO

$$f(\mathbf{w}) = \frac{1}{2n} \sum_{j=1}^n (\mathbf{w}^T \mathbf{x}_j - y_j)^2 + \gamma \|\mathbf{w}\|_1$$



Neural Networks



Conclusions & Future Work

- GIANT's theory beats the existing works.
 - Assume the objective function is **strongly convex** and **Lipschitz smooth**.

Counter-examples


LASSO

$$f(\mathbf{w}) = \frac{1}{2n} \sum_{j=1}^n (\mathbf{w}^T \mathbf{x}_j - y_j)^2 + \gamma \|\mathbf{w}\|_1$$


Neural Networks

Extensions of GIANT (our future work):

Proximal method



Trust-region method



Thank You!